



JHipster 4 Workshop

Introduction

Welcome



- Introduction
- Agenda
- Practical details

Introduction

- Who we are



Introduction

- Who you are



Overview of the workshop

- Goal: being able to set up a complete application with JHipster
 - With a good understanding of the concepts
- Theory (slides) and hands-on labs



Server side frameworks



Spring Boot



Spring Security



Netflix OSS



Consul



Gradle



Maven



Hibernate



Liquibase



MySQL



PostgreSQL



Oracle



MongoDB



cassandra
Cassandra



EhCache



Hazelcast



ElasticSearch



Kafka



Swagger



docker
Docker



Kubernetes



Stack
ELK Stack



Thymeleaf



Gatling



Cucumber

Client side frameworks



HTML5



CSS3



Bootstrap



AngularJS



Angular



jQuery



Websockets



Yarn



Webpack



Bower



Gulp



Sass



Browsersync



Karma



Protractor

Versions of the frameworks & tools

- This workshop is based on the following versions
 - JHipster 4.x
 - Spring Boot 1.5.x with Spring Data
 - Yeoman 1.8.x
 - Webpack 2.2.x
 - Angular 2



Agenda

- Angular 2 Intro
- Introduction to JHipster
- Basic technologies used
- The main generator
- The project structure
- The “entity” sub-generator
- Available development workflows
- Database access
- Security
- Testing
- Deploying to production
- Doing microservices



Practical details - The VM

- Warning: setting up your environment can take some time
 - You should do this before the workshop
 - You need a good network connection
- JHipster can be installed on your machine (better but more complicated)
 - Follow the instructions on <https://jhipster.github.io/installation.html>
 - We will use MySQL during the workshop, so you need to install it from <http://dev.mysql.com/downloads/>
- JHipster can be used within a Virtual Machine (easier to setup)
 - Install VirtualBox <https://www.virtualbox.org/>
 - Install Vagrant from <https://www.vagrantup.com/>
 - Follow the setup instructions from <https://github.com/jhipster/jhipster-devbox>
 - Use the “Quick setup”, using the distribution available on Atlas
 - Modify your VirtualBox configuration to have as much RAM and CPU as possible
 - Change your keyboard settings if you need to

Practical details - checking your development machine

- Do you have an IDE working?
 - ➔ If you use the JHipster “devbox” Virtual Machine, you have Spring Tool Suite installed, just click on the “STS” link on your Desktop
- Do you have MySQL installed?
- Does “yo doctor” work?
- Does “java -version” give you a Java version $\geq 1.8.0$?



Introduction to JHipster

IPPOO



What is JHipster

- An application generator
 - Spring Boot + Angular + tooling
 - 100% free and Open Source
 - Available on Github at <https://github.com/jhipster/generator-jhipster>
- ➔ Come and join us!

What is it good for?

- Business applications
 - Initial focus
 - Single Web Page Application
 - Form-based, CRUD applications
- Web applications
 - More and more used for B2C Webapps
 - Great UI and back-end performance make it a compelling solution
- Mobile applications
 - Responsive UI
 - With a native UI accessing the REST API

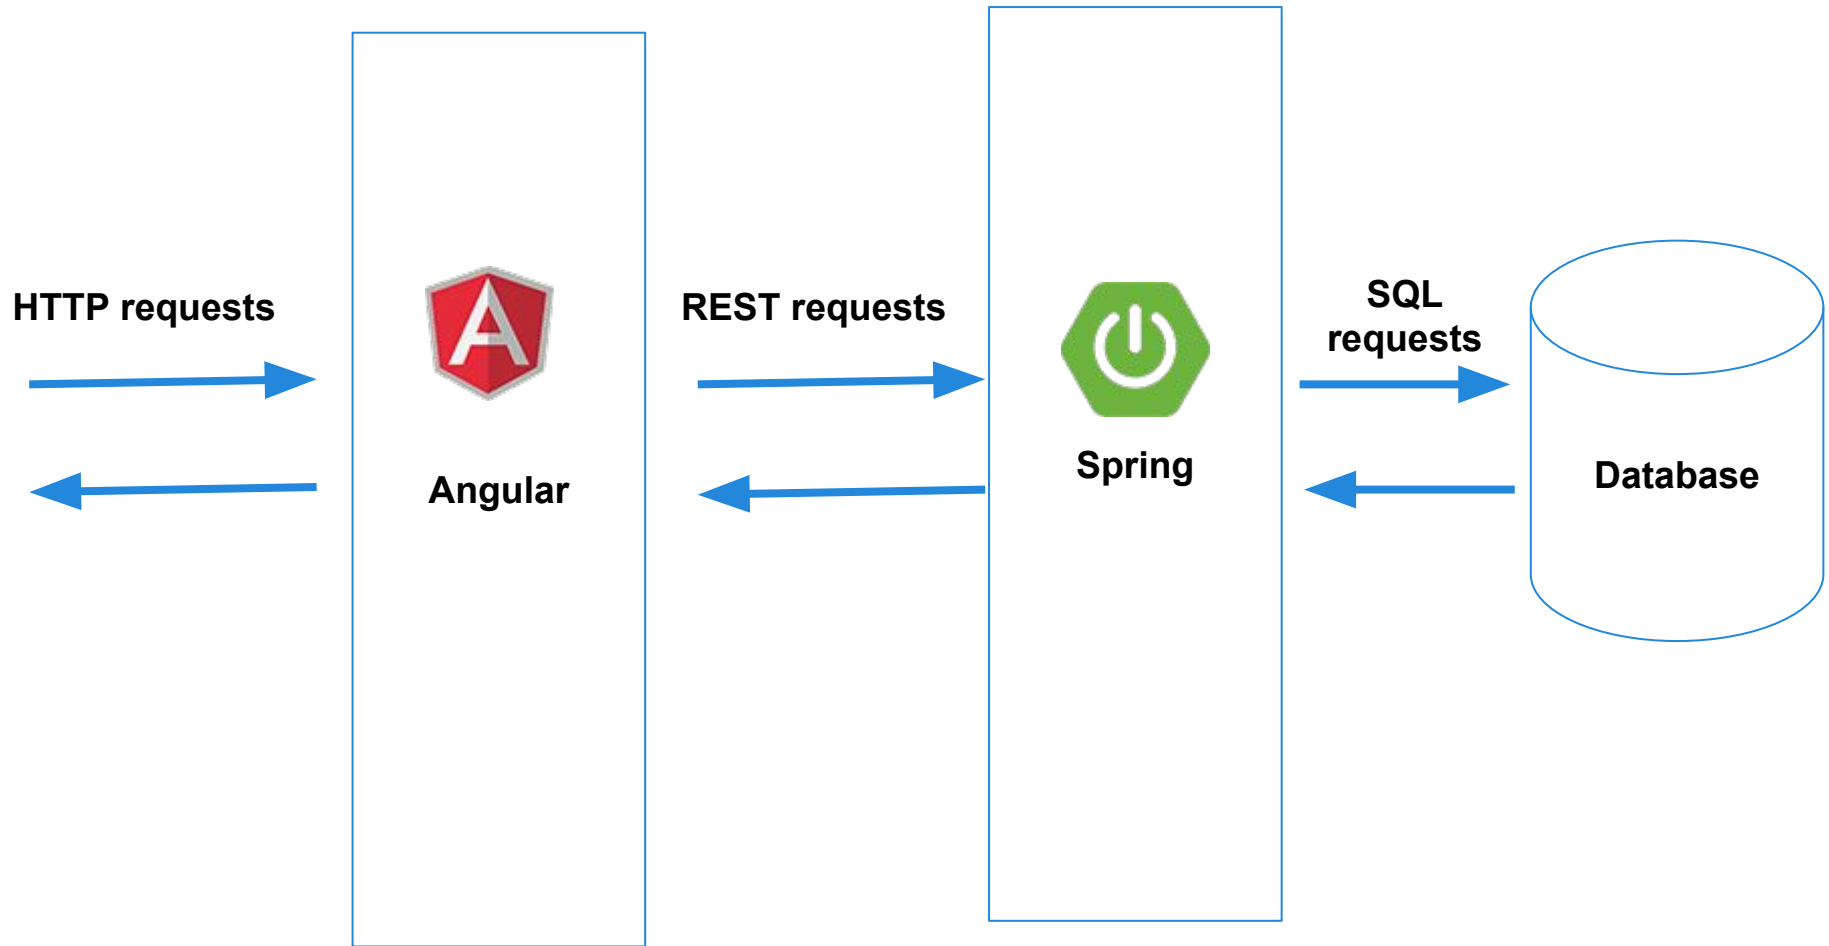




Architecture

- The classical JHipster architecture is that of a Single Page Application
 - ➔ Your database may vary (supports SQL, MongoDB, Cassandra)
 - ➔ Your services may vary (Elasticsearch, caching...)
- A main, static HTML page is loaded
 - ➔ Launches an Angular application
 - ➔ This application does REST requests to the back-end
 - ➔ The back-end handles those requests, add business services and security
 - ➔ And data is stored in a database

Architecture



What is being generated?

- The application skeleton
 - Spring Boot on the back
 - Angular on the front
 - Tooling is configured so you can get started quickly
- You can also use “sub-generators” to add more features
 - The main one is the “entity sub-generator” that generates CRUD interfaces





Configured tooling

- JHipster comes packed with tools to ease your development workflow
 - ➔ The main goal of this masterclass is to learn how to use them efficiently
- Server-side tools
 - ➔ Spring Boot adds a lot of features, including “hot reload”!
 - ➔ Liquibase for database updates
- Client-side tools
 - ➔ Yarn, Webpack & BrowserSync for “live reload” of your front-end application
- Many “dotfiles” for your IDE or Git
- Docker files for running services in a container

Generating your application graphically

- JHipster builds on top of Yeoman
 - A question/answer system to generate your application
 - Very powerful, but can be annoying if you have a lot of things to generate
- Applications can be generated from a standard .xmi file (coming from your own UML editor)
 - You develop your model graphically using an UML editor, and export it as an .xmi file
 - The JHipster UML project is then able to generate the entities automatically
 - More information at: <http://jhipster.github.io/jhipster-uml/>
- The free, online “JDL Studio” can be used to generate the application using a domain-specific language
 - <http://jhipster.github.io/jdl-studio/>
 - Download the generated file and run it with JHipster

JHipster best practices

- Follow each underlying technology “best practices”
 - Angular best practices, Spring best practices, etc...
 - Only change them if there is a good reason to
- Use the provided library versions
 - It's hard work to have them all working together
- Follow the workflows provided by JHipster
 - They are here to help you
 - There is usually a very good reason to use them in the recommended way
- You have a great working environment out-of-the-box, don't break it!
 - Front-end and back-end updates should be automatic and fast
 - Production deployment should be easy
- Git is your friend
 - Every mistake (including in the database!) should be easy to rollback with Git





Basic technologies used

IPPOO

Spring Boot

- Hugely popular, on top of the hugely popular Spring Framework
- Spring Boot helps configuring and getting started with Spring
 - ➔ Convention over configuration
 - ➔ Automatic configuration, with clean and documented ways to adapt it to your needs
- JHipster doesn't get in the way of Spring Boot
 - ➔ Configures it, so that it works great with Angular
 - ➔ Builds on top of it and adds new features: mainly for management and security
 - ➔ Very scarcely, corrects some Spring Boot issues, but tries as much as possible to give those back to Spring Boot

Spring Data

- One of the many Spring sub-projects, integrated in a Spring Boot starter
- Greatly simplifies the way JPA works with a dedicated module
 - ➔ Automatic query generation, based on interface method names
- In a typical JHipster application, you will need to create those queries yourself
 - ➔ It's important to understand Hibernate/JPA
 - ➔ At first Spring Data JPA looks a bit complex, but when you are used to it, it's really quick and efficient
- If you use a NoSQL database
 - ➔ MongoDB: Spring Data MongoDB works the same (there even are portable methods, which work for both technologies!)
 - ➔ Cassandra: it was deliberately chosen not to use Spring Data Cassandra, mainly for performance and stability reasons



Spring MVC REST

- The de-facto standard REST framework in Java
 - ➔ Great performance, lots of documentation on the Internet
 - ➔ Very close to most other REST frameworks anyway
- JHipster creates “CRUD” operations automatically
 - ➔ GET, PUT, POST, DELETE verbs from HTTP
 - ➔ JSON objects with Jackson

Spring Security

- One of the oldest Security frameworks in Java
 - Widely used
 - Lots of plugins and features
- Secures both URLs and Java methods
 - Use both features together for extra security!
- JHipster adds a lot of features on top of it
 - Several different authentication schemes
 - Great UI with Angular
 - Complete database schema
 - Added features, such as invalidation of the persistent tokens



Angular

- Hugely popular JavaScript framework
- Easy to use at first, but hard to have a good modular design
- JHipster provides the architecture and many building blocks
 - ➔ Clean and modular design, which allows your application to grow without trouble
 - ➔ REST clients that connect to what Spring MVC REST provides
 - ➔ Built-in support for Spring Security
 - ➔ i18n, both on the client and on the server side
- By default, JHipster generates a Single Page Application that provides a rich and fluid user experience

Bootstrap

- Hugely popular JS/CSS framework
- JHipster uses its classes extensively
 - ➔ Responsive UI out-of-the-box
 - ➔ Tables, alerts, popups, menus... all come from Bootstrap
- Great for Web Designers
 - ➔ As JHipster uses Bootstrap correctly, it's easy to add a new theme on top of a JHipster application
 - ➔ Modify the provided theme or install a new one

Yeoman

- The tool that is used to create JHipster
- As a JHipster user, you will mostly see it when you call “yo”
- Provides many interesting features for advanced users
 - Configuration files are easy to modify
 - Re-generating a project is seamless
 - Templates are great (and have been improved by the JHipster team!)



Maven & Gradle

- Your build tools for the Java world
- JHipster provides wrappers for both, so you don't have to install them
 - ➔ Best practice: each build tool has its own version, to guarantee an immutable build
 - ➔ Of course you can still install them if you want to!
- JHipster provides libraries and plugins versions, all configured
 - ➔ Extends Spring Boot, which also provides those versions
 - ➔ Only modifies a version on top of Spring Boot when there is a good reason to do so
 - ➔ Think carefully before upgrading a version, as both the Spring Boot team and the JHipster team have validated it!
- Public repositories can be searched for installing/updating libraries
 - ➔ <http://search.maven.org/>
 - ➔ <https://repository.sonatype.org/>



The main generator

IPPOO



Available options

- JHipster provides many options when creating an application
 - It can generate an application tailored to your specific needs
 - When in doubt, use the default options
- The most important options are:
 - Client framework: AngularJS or Angular
 - Type of authentication
 - Database to be used in production and in development
 - Hibernate cache
 - Usage of Elasticsearch
 - Build tools that should be configured: Maven/Gradle
 - i18n
 - Testing frameworks

Generated files

- JHipster generates a lot of files
 - ➔ The node_modules directory is used internally by NPM, you should ignore it
 - ➔ Beside this, it is a classical Maven project structure
 - With "src/main/java" and "target" directories, as usual
 - ➔ In the "src/main/webapp" folder you will find the whole Web application, including the Angular 2 application
- The Spring Boot application has the usual Java package hierarchy
 - ➔ "domain" for your domain objects, "repository" for your repositories, "web/rest" for you REST endpoints
 - ➔ Layers are separated by technical concerns, and not by business concerns
- The Angular application is separated in modules
 - ➔ Follows the Angular style guide
 - ➔ <https://angular.io/docs/ts/latest/guide/style-guide.html>



The .yo-rc.json file

- This is the internal Yeoman configuration file
 - ➔ You can tweak it manually if you want to re-generate your application
 - ➔ You should send it to the team if you have an issue
- JHipster also has its own configuration files in the .jhipster directory
 - ➔ At the moment, this is only for generating entities
 - ➔ This is what JHipster UML generates

How to start correctly a new project

- Create an empty directory in which you will work
- Generate the application
- Test if everything is OK
 - ➔ run “mvn test” for the back-end
 - ➔ run “yarn test” for the front-end
- Commit it to Git
 - ➔ JHipster generates a .gitignore file for you, so you can commit your files straight after generation
- Import it into your IDE
 - ➔ Spring Tool Suite, Eclipse, IntelliJ IDEA, should all work
 - ➔ More information on your IDE configuration at http://jhipster.github.io/configuring_ide.html

How to update an existing project

- If you run again “yo jhipster” or “yo jhipster:entity” on an existing project, JHipster will re-use its configuration files to re-generate the project
 - ➔ If you have upgraded to a new JHipster version, you will get newer versions of your files
- The trouble when updating an existing project is that JHipster will overwrite all files
 - ➔ Yeoman provides a mechanism for this, that will warn you and allow you to make a diff to see what’s being changed
 - ➔ But for big updates, this isn’t enough
- The best solution is to use Git and the Git tools to do the merging
 - ➔ We recommend using a graphical tool, like SourceTree:
<https://www.sourcetreeapp.com/>



Lab 1: setting up a project

IPPoon

The sample application

- The sample application is a Bug Tracker
- It can handle several projects
- Each project has several tickets
- Tickets can have labels, like “bug” or “enhancement”
- Tickets are assigned to users



Your first JHipster project (1/2)

- Create a new directory called “BugTracker”
- Run “yo jhipster” in that directory
 - When the database is selected, choose a MySQL in development and in production
 - When the client framework is selected, choose Angular 2.x
 - Do NOT enable internationalization support
 - Keep all other options as default
- Validate your project
 - “./mvnw test” for the server side
 - “yarn test” for the client side
- Open up your application on your IDE
 - If you use the JHipster “devbox”, the shortcut is on your desktop
 - Configure your IDE in order to import the project

Your first JHipster project (2/2)

- Run MySQL
 - ➔ The recommended way is to use Docker:
`docker-compose -f src/main/docker/mysql.yml up`
 - ➔ You can also install MySQL manually:
`apt-get install mysql`
- Create a MySQL schema
 - ➔ The schema name is by default your application's name
- Configure your "application-dev.yml" to point to your MySQL schema
- Run your application
 - ➔ Run the "Application" class
 - ➔ Go to your application at <http://127.0.0.1:8080>
 - ➔ Check your database schema



Project structure

IPPOO

Java project structure

- JHipster generates a “standard” Maven project
 - Usual Maven project structure
- Java packages match the usual technical layers
 - Configuration
 - Repositories
 - Domain objects
 - Services
 - REST endpoints



Configuration files

- Configuration files are generated in the “src/main/resources/config” directory
- Spring Boot is configured using the “application.yml” files
 - Uses the YAML format
 - Fully documented (and your IDE might even be able to autocomplete them!)
 - See <http://docs.spring.io/spring-boot/docs/current/reference/html/common-application-properties.html>
 - One base “application.yml” file, and then one file per Spring profile
 - Each key can be overridden using an environment variable
- Liquibase is configured in the “src/main/resources/config/liquibase” subdirectory
 - One master file that points to several changelogs
 - Automatically updated when running “yo jhipster:entity”
 - Will require editing when running “mvn liquibase:diff”



Web project structure

- The Web project is in the “src/main/webapp” folder
 - ➔ Contains many best practices for Web development
- “app” contains the Angular application
- “content” contains static resources, such as images and CSS files
- Note that a “target/www” folder will be created during the project’s production packaging, and will contain a minified version of the web project



Angular project structure

- “app” contains the main application
- “app/account” contains the user account management UI
- “app/admin” contains the administration UI
- “app/blocks” contains common building blocks like configuration and interceptors
- “app/home” contains the home page
- “app/shared” contains common services like authentication and internationalization
- “app/entities” is where your entities will get generated
- “app/layouts” contains UI code such as the navigation bar and error pages



The “entity” sub-generator

IPPOO

Generating a simple CRUD entity

- The “entity” sub-generator can be run in the application folder, to add entities to the application
 - ➔ Run “yo hipster:entity Ticket” to generate the “Ticket” entity
- It generates a “domain object”, as well as all the required infrastructure, back-end and front-end code
 - ➔ Liquibase changelog
 - ➔ JPA entity
 - ➔ JPA repository
 - ➔ Spring MVC REST controller with GET, PUT, POST, DELETE methods
 - ➔ Angular and HTML code
 - ➔ Tests
- This generation depends on your application’s main configuration
 - ➔ Supports SQL, Cassandra, MongoDB
 - ➔ Supports Elasticsearch...

Field types

- Many field types are provided out of the box
 - For each field, you will be asked to define its type, validation rules, pagination...
 - Each type is configured to be used in HTML, in Java, and in the database
- Available types are:
 - String
 - Numeric types: Integer, Long, Float, Double, BigDecimal
 - Date types (from the Java 8 Date/Time API): LocalDate, ZonedDateTime
 - Boolean
 - Enum (which generates a Java enum type)
 - Blob, which can handle both a generic binary file or an image
- If you need a type that doesn't exist yet, use the closest available type (or use "String") and then do a refactoring with your IDE

Validation

- For each field type, validation is available
- On the client-side, the validation system from Angular 2 is used
 - ➔ Works with Bootstrap to provide a nice UI
- On the server-side, Bean Validation works with Spring MVC REST
 - ➔ Bean validation is executed automatically by Spring MVC REST when the @Valid annotation is present at the controller level
- We try to have a coherent validation between the client and the server side
 - ➔ Not always possible, by example for Regexp patterns
 - ➔ Forces JHipster to use the “lowest common denominator”
 - ➔ For advanced usage, you will need to tweak the generated code manually

Relationship types

- Relationships only work if you have selected an SQL database
- JHipster supports relationships generation between entities
 - one-to-one
 - one-to-many
 - many-to-one
 - many-to-many
- Those relationships can be bidirectional or not
- Relationships rely on database (and JPA) relationships
 - generates the Liquibase and JPA code
 - also generates specific Angular 2/HTML front-end code
- Warning! Don't forget JPA relationships are lazy by default
 - JHipster uses the underlying technology's default configuration
 - Use eager relationships or transactions to solve this common issue with JPA



Entity creation order

- Relationships force the creation of some entities before others
- This comes from the database
 - ➔ You can't add a foreign key pointing to a table that doesn't exist yet
- For a one-to-many relationship
 - ➔ You must first do the "one" side (as the foreign key is on the "many" side)
- For a many-to-many relationship
 - ➔ You must first build the entity that doesn't own the relationship



The “User” entity

- The “User” entity is a specific entity, as it is already created by JHipster
 - ➔ It is used by JHipster for authentication (with Spring Security)
 - ➔ It can be modified as any generated Java code
 - ➔ There is an administration UI specifically provided out of the box for this entity, to manage users
- A “many-to-one” relationship can be added to the “User” entity
 - ➔ Example: several tickets are assigned to one user

Pagination options

- Pagination is optional
- On the server-side, a “GET” request to get all objects has by default a pagination of 20 objects
 - Can be configured through Spring MVC REST/Spring Data
- On the client-side, 3 different options are available
 - Simple pager
 - Pagination links
 - Infinite scrolling





Lab 2: generating entities

IPPOON

Entity 1: the “Project”

- Create an entity called “Project”: `yo jhipster:entity project`
- It has only one field, “name”, of type “String”
- Hit “compile” on your IDE, everything should be automatically updated
- Run your application
 - ➔ Either from your IDE like in the previous lab, or type “./mvnw”
 - ➔ Run Browsersync by typing “yarn start”, this will open up a Web browser using <http://localhost:9000>

Entity 1 solution



.hipster/Project.json

```
{
  "fluentMethods": true,
  "relationships": [],
  "fields": [
    {
      "fieldName": "name",
      "fieldType": "String"
    }
  ],
  "changelogDate": "20151118144708",
  "dto": "no",
  "service": "no",
  "entityTableName": "project",
  "pagination": "no"
}
```

Entity 2: the “Label”

- Create an entity called “Label”: `yo jhipster:entity label`
- It has only one field, “label”
 - it is of type “String”
 - it must be validated: it is a required field, with a minimum length of 3
- It has a relationship to the entity “Ticket”
 - It is a many-to-many relationship
 - The label is the non-owning side
- Hit “compile” on your IDE, **it is going to FAIL** as the “Ticket” entity doesn’t exist yet

Entity 2 solution

.jhipster/Label.json



```
{
  "fluentMethods": true,
  "relationships": [
    {
      "relationshipName": "ticket",
      "otherEntityName": "ticket",
      "relationshipType": "many-to-many",
      "ownerSide": false,
      "otherEntityRelationshipName": "label"
    }
  ],
  "fields": [
    {
      "fieldName": "label",
      "fieldType": "String",
      "fieldValidateRules": [
        "required",
        "minlength"
      ],
      "fieldValidateRulesMinlength": "3"
    }
  ],
  "changelogDate": "20151118145609",
  "dto": "no",
  "service": "no",
  "entityTableName": "label",
  "pagination": "no"
}
```

Entity 3: the “Ticket”

- Create an entity called “Ticket”
 - If you make a mistake during those questions/answers, you will need to restart the entity creation
- It has several fields
 - “title”, of type String, which is required
 - “description”, of type String
 - “dueDate”, of type LocalDate
 - “done”, of type Boolean
- It has several relationships
 - with “Project”, it has a many-to-one relationship, using the Project “name” to display the relationship
 - with “User”, it has a many-to-one relationship named “assignedTo”, using the User “login” to display the relationship
 - The “User” entity is generated by default by JHipster
 - with “Label”, it has a many-to-many relationship, which it owns, using the Label “label” to display the relationship
- It uses pagination links

Entity 3 solution

```
{
  "fluentMethods": true,
  "relationships": [
    {
      "relationshipName": "project",
      "otherEntityName": "project",
      "relationshipType": "many-to-one",
      "otherEntityField": "name"
    },
    {
      "relationshipName": "assignedTo",
      "otherEntityName": "user",
      "relationshipType": "many-to-one",
      "otherEntityField": "login",
      "ownerSide": true,
      "otherEntityRelationshipName": "ticket"
    },
    {
      "relationshipName": "label",
      "otherEntityName": "label",
      "relationshipType": "many-to-many",
      "otherEntityField": "label",
      "ownerSide": true,
      "otherEntityRelationshipName": "ticket"
    }
  ],
}
```

```
"fields": [
  {
    "fieldName": "title",
    "fieldType": "String",
    "fieldValidateRules": [
      "required"
    ]
  },
  {
    "fieldName": "description",
    "fieldType": "String"
  },
  {
    "fieldName": "dueDate",
    "fieldType": "LocalDate"
  },
  {
    "fieldName": "done",
    "fieldType": "Boolean"
  }
],
"changelogDate": "20151118154200",
"dto": "no",
"service": "no",
"entityTableName": "ticket",
"pagination": "pagination"
}
```

[jhipster/Ticket.json](#)



Available development workflows

IPPOO

Running the Java application

- The Java application can be run in three ways
 - mvn
 - run the “Application” class
 - run the application as a “Spring Boot application” if your IDE supports Spring Boot
- It can be run with 2 profiles
 - A profile is a Spring Profile as well as a Maven Profile (JHipster generates those consistently)
 - The “dev” profile is the one selected by default, and is used for development
 - The “prod” profile is for production use
- The “dev” profile has Spring Boot devtools enabled
 - Uses a specific classloader to restart the application faster
 - This is a must-have feature, you have to use it!



Using Maven (or Gradle)

- “mvn” runs the default “mvn spring-boot:run” goal
 - Runs the application
- “mvn clean” will cleans the application’s “target” directory
 - Cleans up Elasticsearch indexes and H2 disk-based persistence database
- “mvn compile” compiles the classes
 - Works with Spring Boot devtools if you have “mvn” running in another terminal
- “mvn test” runs the tests
 - Unit tests and integration tests



Using Yarn

- “yarn start” runs the browser on port 9000 by default
 - ➔ Automatically watch files
 - ➔ Runs the front-end application with Browsersync
- “yarn test” runs the front-end tests
 - ➔ With karma.js and phantom.js

Working with BrowserSync

- Browsersync is a great tool for Web development
 - ➔ <http://www.browsersync.io/>
- It synchronizes different browsers
 - ➔ For testing different browsers or resolutions at the same time
 - ➔ Also works from external devices: control one or several phones automatically
- Synchronizes your behavior: clicks, scrolls, form inputs...
- Refreshes the application when a file changes
 - ➔ HTML, CSS, and JavaScript
 - ➔ No more manual browser refresh!

Database updates with Liquibase

- Liquibase versions databases changes
 - Very handy to update and maintain a database's structure
 - Allows to have several people work on the same database together
- JHipster generates Liquibase changelogs for you
 - When the main generator is run
 - When the entity sub-generator is run
- JHipster also configures the Maven Liquibase Hibernate plugin according to your current database state
 - The plugin will connect to your database and generate a changelog between your current JPA configuration and your database schema
 - Works well with MySQL or Postgresql database (support for disk-based H2 too)
 - Edit a JPA entity, compile it, and run `"mvn liquibase:diff"` to have a new changelog in `"src/main/resources/config/liquibase/changelog"`. You will need to include it in the `"src/main/resources/config/liquibase/master.xml"` file.



Lab 3: common workflows

IPPOon

Workflow 1: updating client-side code

- Go to your application's tickets page at <http://localhost:9000/#/tickets>
- With your IDE, look for the ticket.component.html page
 - Remove the "Description" column
 - Change the format of the "DueDate" to "MM/dd/yyyy"
 - Each time you change this file, your browser should refresh automatically



Workflow 2: Liquibase

- Add a new field to the “Ticket” entity
 - A ticket has a “priority” which is an integer
 - Add this field to the Ticket entity, and generate its getter/setter methods
 - Compile the Ticket class
 - Open up your “pom.xml” file and configure your Liquibase plugin to be sure it points to your MySQL database
 - Check your schema name, login and password
 - Run “mvn liquibase:diff”
 - Have a look at the generated changelog, and add it to the master.xml
 - Compile the application again to make Liquibase execute the changelog
 - Use MySQL workbench to check that the field has been added to the database
- (optional) modify the user interface to see the priority field
 - In the “ticket-dialog.component.html” file, add a new text input (re-use description)
 - In the “ticket.component.html” file, add a new “priority” column which will display the priority



Database access

IPPOO

Spring Data JPA

- Spring Data JPA provides many helpers methods to ease working with JPA/Hibernate
 - ➔ Simple CRUD repositories are automatically provided
 - ➔ More complex queries can be generated dynamically using method names
 - ➔ Provides pagination and ordering API
- For example, to have all the “Foos” ordered by the “bar” field:
 - ➔ `fooRepository.findAllOrderByBarDesc`
- More information on the Spring Data JPA project at <http://projects.spring.io/spring-data-jpa/>

Hibernate 2nd level cache

- Using a 2nd level cache is the top priority if you want good performance using JPA
- JHipster provides 3 options
 - ➔ No cache
 - ➔ Ehcache (the default)
 - ➔ Hazelcast (a distributed cache)
- Caches will be generated using default values, which are quite low for production
 - ➔ You need to tune them, depending on your usage and available RAM
 - ➔ Ehcache is monitored, and can be observed in the application's administration/monitoring screen: use it to tune your cache configuration accordingly

MongoDB

- MongoDB is a NoSQL database, mainly used to store JSON documents
 - ➔ More information at <https://www.mongodb.org/>
- JHipster supports MongoDB
 - ➔ Using Spring Data MongoDB
 - ➔ Using Mongeez for database upgrades (instead of Liquibase)



Cassandra

- Cassandra is a NoSQL database providing high scalability and availability
 - ➔ More information at <http://cassandra.apache.org/>
- JHipster supports Cassandra
 - ➔ Using directly the DataStax Java Driver for Cassandra



Lab 4: database access

IPPON



Working with Spring Data JPA

- We want to have the most urgent tickets on top of the list
 - ➔ Currently they are not ordered
- Modify the TicketResource REST controller to use a new Spring Data repository method
 - ➔ Find the right method name to have all the entities, ordered by due date (descending)
 - ➔ Use your IDE to generate this method in the repository
 - ➔ Thanks to the Spring Boot devtools, your application should restart automatically (and very quickly!) with your new repository method



Solution

TicketResource.java

Page<Ticket> page =

ticketRepository.findAllOrderByDueDateDesc(pageable)



Security

IPPOO

Available options

- For authentication and authorization, JHipster configures Spring Security, the corresponding database schema and user interface
- 4 options are currently available:
 - ➔ Session-based authentication
 - ➔ Session-based authentication with social login enabled
 - ➔ OAuth2
 - ➔ JWT
- OAuth2 and JWT are stateless mechanisms, and can be easier to scale on a cluster



Remember-me

- “Remember-me” is a Spring Security option that allows to a user to stay connected for 30 days
 - ➔ Works for session-based and JWT authentications
- For session-based authentication, JHipster adds many features on top of what is provided by Spring Security
 - ➔ Tokens are far more secure than the default Spring Security implementation
 - ➔ Database table to store the user’s tokens
 - ➔ User interface in Angular 2, including support for invalidating current tokens

CSRF protection

- CSRF attacks use a forged link to your application in order to execute commands, using your current authentication
 - ➔ With JHipster, this could happen if you have session-based authentication: the attacker would use your session cookie to be authenticated
- JHipster uses Spring Security's CSRF protection to avoid this issue
 - ➔ A token is sent in the browser's headers, and must be validated
 - ➔ Requests using the HTTP GET verb are not validated, as they are not modifying anything
 - ➔ Angular 2 is specifically configured to understand the same token as JHipster



Using the “User” entity

- The User entity is a specific entity generated by JHipster
 - ➔ It represents a “User” domain object
 - ➔ It is used by Spring Security to check the user’s authentication and authorizations
- It can be used throughout your application
 - ➔ More fields can be added
- A many-to-one relationship to the “User” can be created
 - ➔ A specific Spring Data JPA query is generated so the other entity can be filtered on the current user (a common use case)



Lab 5: security

IPPOon

Securing methods

- Only the “admin” users should have the permission to delete existing tickets
 - ➔ Add a security annotation on the TicketResource REST controller for users with the “ROLE_ADMIN” permission only can delete the resource
 - ➔ Hit “compile”, your application should restart automatically
 - ➔ Log in the application with “user/user”, and try to delete a ticket: you should have an HTTP 403 (forbidden) error
- (Optional) Remove the “delete” button when the user doesn’t have the ROLE_ADMIN permission
 - ➔ Hint: use “!hasAnyAuthority”

Solution

TicketResource.java

```
@Secured(AuthoritiesConstants.ADMIN)
public ResponseEntity<Void> deleteTicket(...)
```

ticket.component.html

```
<button type="submit"
        [routerLink]="['/', { outlets: { popup: 'ticket/' + ticket.id + '/delete' } }]"
        replaceUrl="true"
        class="btn btn-danger btn-sm"
        jhiHasAnyAuthority="ROLE_ADMIN" >
```




Testing

IPPOO

Testing with Spring

- On the server side, JHipster provides
 - ➔ Unit tests using vanilla JUnit
 - ➔ Integration tests using Spring's test context framework
- Integration tests provide
 - ➔ An in-memory database, with the Liquibase changelogs applied
 - ➔ Tooling to test the application's REST endpoints
- Those tests are located in "src/test/java"
 - ➔ They can be run inside an IDE
 - ➔ "mvn test" runs all the Java tests



UI testing with Karma.js

- Karma.js is a framework to do unit tests in JavaScript
 - ➔ Uses mocks to simulate the back-end
 - ➔ Uses Phantom.js to run a browser from the command line
- Those tests are located in “src/test/javascript”
 - ➔ “yarn test” runs all the JavaScript tests

Performance testing with Gatling

- Gatling is used to do performance tests on the back-end
 - Can simulate a high number of concurrent users
 - Provides a Scala DSL to code the test scenarios
- Those tests are located in “src/test/gatling”
 - “mvn gatling:execute” allows to select which test to run
 - A report is generated at the end of the test, with detailed performance results





Lab 6: testing

IPPoon



Testing on the Java side

- A ticket's priority should not be negative: add a "@Min(0)" Bean Validation annotation on the "priority" field we have created
- Test this new validation rule
 - ➔ Add a new method to TicketResourceIntTest
 - ➔ This method should test that if the priority is "-1", a "400 Bad Request" is returned by the server
 - ➔ (optional) Also test that a priority of "1" returns a "201 Created"

Solution

Ticket.java

@Min(0)

private int priority;

TicketResourceIntTest.java

@Test

@Transactional

```
public void checkPriorityIsNotNegative() throws Exception {
    int databaseSizeBeforeTest = ticketRepository.findAll().size();
    ticket.setPriority(-1);
    restTicketMockMvc.perform(post("/api/tickets")
        .contentType(TestUtil.APPLICATION_JSON_UTF8)
        .content(TestUtil.convertObjectToJsonBytes(ticket)))
        .andExpect(status().isBadRequest());

    List<Ticket> tickets = ticketRepository.findAll();
    assertThat(tickets).hasSize(databaseSizeBeforeTest);
}
```



Deploying to production

IPPOO

The “prod” profile

- JHipster has a “prod” profile, that packages an application for production use
 - ➔ It will optimize the front-end for production
 - ➔ It will tune the back-end for production, consistent with the front-end
- It is run with Maven
 - ➔ “mvn package -Pprod” will generate a production WAR file in the “target/” directory
 - ➔ “mvn -Pprod” will run directly the application in “prod” mode
- The generated WAR file contains the production code
 - ➔ To run it with the Spring “prod” profile, use the standard Spring Boot configuration:
“./jhipster-0.0.1-SNAPSHOT.war --spring.profiles.active=prod”




Front-end minification

- JHipster will use Yarn to minify the application front-end
- All JS/CSS/HTML files will be minified, and the result will appear in the “target/www” folder
 - ➔ You can compare the “src/main/webapp/index.html” file used in development with the “target/www/index.html” file used in production
 - ➔ Minified files will have a checksum in front of their name, so they can be cached effectively
 - ➔ You can use the Chrome Dev Tools, including the audit tool, to check that everything has been minified

Server-side tuning

- In production mode, JHipster optimizes the web application
 - ➔ HTTP requests are GZipped (if they are >1kb)
 - ➔ Cache headers are added to all static files (see the previous slide about the checksum added in front of each file)
- The application uses the “application-prod.yml” configuration file
 - ➔ Including the specific configuration properties for using the production database
- Spring Boot devtools are disabled



Deploying the WAR file

- The result of the “production build” is a WAR file
- It is an executable WAR file, that contains an embedded Tomcat server
 - ➔ It can be run in production with
`“./jhipster-0.0.1-SNAPSHOT.war --spring.profiles.active=prod”`
 - ➔ It can be deployed in a Java EE application server that supports the Servlet 3.0 specification, using the “spring.profiles.active=prod” environment variable

Deploying to a cloud provider

- JHipster provides several sub-generators to deploy the application to the cloud
 - ➔ Cloud Foundry
 - ➔ Heroku
 - ➔ Amazon Web Services
- Each of those providers have specific prices and features
- If you selected “MySQL” and “Elasticsearch” options when generating your application, your cloud provider must also provide options for both



Deploying to the Cloud Demo

IPPOon

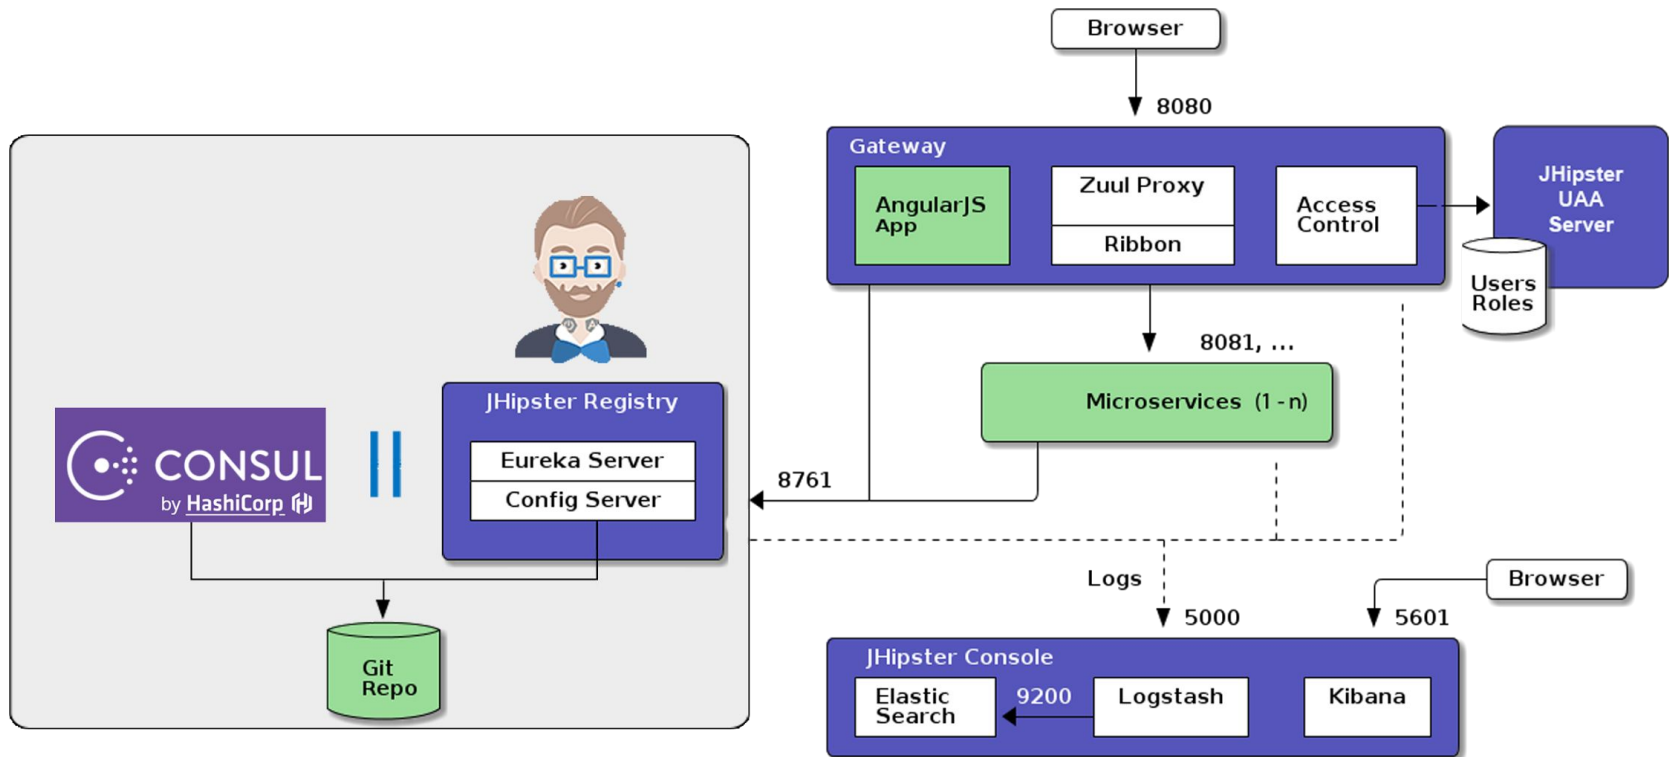


Doing microservices

IPPOO

The big picture

NETFLIX | OSS +  +  docker



 elastic +  logstash +  kibana



A microservice with JHipster

- Normal JHipster application, focused on the back-end
 - ➔ All databases work as expected
 - ➔ Main JHipster options work as usual
 - ➔ Sub-generators work
- No graphical user interface (no Angular)
- No user management code



A gateway with JHipster

- Normal JHipster application, focused on the front-end
 - Can also have a back-end: it has a database, and can have full-fledged entities
 - All usual options work as usual
 - Sub-generators work
- But it is also a “router” to microservices
 - Does load balancing and circuit breaking
 - Very high performance
 - Handles security and user management for microservices
 - Optionally, can have rate limiting features
- And it has a specific UI
 - Can generate front-end code (Angular) based on an existing microservice



The JHipster Registry

- Runtime component provided by JHipster
- Fully Open Source (Apache 2 license)
- Service Registry based on Spring Cloud Eureka
 - ➔ All services register themselves on the JHipster Registry
 - ➔ Allows load balancing on the gateways
 - ➔ Allows microservice scalability and cluster configuration
- Configuration server based on Spring Cloud Config
 - ➔ Sends configuration data to all services
 - ➔ Useful to version, tag, rollback configurations
 - ➔ Allows to store “sensitive” information like database passwords
- NEW: Can choose to use Hashicorp’s Consul instead of the JHipster Registry



JHipster Console

- Monitoring console based on the ELK stack
 - Elasticsearch, Logstash, Kibana
 - Aggregates logs from microservices and gateways
 - Provides pre-defined dashboards
- Logs are sent by each JHipster application
 - Log messages sent by using the logback API: "log.debug()"
 - Dropwizard Metrics data dumped regularly to the logs, with detailed information from the JVM, Spring Beans, etc.
- Alerting is also available
 - Using Elastalert from Yelp

Running everything with Docker Compose

- JHipster applications are always configured to run with Docker
 - ➔ Services like databases or Elasticsearch have a pre-configured Docker Compose configuration in the “src/main/docker” directory
 - ➔ Applications can be packaged as Docker images using Maven or Gradle
 - ➔ Run “mvn -Pprod package docker:build” to generate an image of your current application in production
 - Uses Alpine Linux and Open JRE
 - Very light and secured
- When doing microservices, JHipster provides the “docker-compose” sub-generator
 - ➔ Allows us to have a Docker Compose configuration of everything, so running the whole architecture is very easy
 - ➔ Includes the JHipster Registry and (optionally) the JHipster Console
 - ➔ Allows us to scale services easily using “docker-compose scale”



Lab 7: creating a simple microservice

IPPON

Generating a microservice

- Create a new directory called “microservices”
- In that directory, create an application called “catalog”
 - ➔ Create a “catalog” sub-directory
 - ➔ Generate a “microservice application” in that directory
 - Choose a MySQL database
 - Choose the (default) HazelCast cache
 - Add an Elasticsearch search engine
 - ➔ Create a “product” entity in that application
 - Add the entity using “yo jhipster:entity Product”
 - A product has a name, description, price...
 - ➔ Run “mvn test” to test your project
- Generate a Docker image of your project
 - ➔ Run “mvn -Pprod package docker:build”

Generating a gateway

- Go to the “microservices” directory
- In that directory, create an application called “gateway”
 - ➔ Create a “gateway” sub-directory
 - ➔ Generate a “microservice gateway” in that directory
 - Choose a MySQL database
 - ➔ Create a “product” entity in that application, that is the front-end to the microservice we have generated
 - Add the entity using “yo jhipster:entity Product”
 - Confirm that you will generate that entity from an existing microservice
 - Write the path to your existing “Catalog” microservice
 - ➔ Run “mvn test” and “yarn test” to test your project
- Generate a Docker image of your project
 - ➔ Run “mvn -Pprod package docker:build”

Building and running with Docker

- Go to the “microservices” directory
- In that directory, create a sub-directory called “docker-compose”
 - Go that directory and run “yo jhipster:docker-compose”
 - Confirm you want to use the parent (“../”) directory
 - Choose to use both the Catalog and the Gateway applications
 - Select ELK monitoring
- Run everything
 - Just type “docker-compose up -d”
- Test
 - Go to <http://127.0.0.1:8080/> to have access to the gateway
 - Create some entities, and test the Elasticsearch is working
 - Go to <http://127.0.0.1:8761/> to have access to the JHipster Registry
 - Go to http://127.0.0.1:5601 to have access to the JHipster Console

Scaling with Docker

- Scale the “catalog” microservice with Docker
 - ➔ Run “docker-compose scale catalog-app=2”
- A second instance of “catalog” is running
 - ➔ As it uses HazelCast, a distributed cache will be automatically configured between both instances
 - ➔ This second instance will be available in the JHipster Registry and in the gateway’s admin screen
 - ➔ It will also be automatically monitored by the JHipster Console
- You can launch new instances, and kill existing ones, to see how the architecture handles failure, circuit breaking and load balancing
- When you have finished, just run “docker-compose down” to destroy everything!



Digital . Technologies . Hosting



PARIS
BORDEAUX
NANTES
LYON
WASHINGTON DC
NEW-YORK
RICHMOND
MELBOURNE
MARRAKECH

contact@ippon.fr
www.ippon.fr - www.ippon-hosting.com - www.ippon-digital.fr
@ippontech
-
01 46 12 48 48

support@stormpath.com
www.stormpath.com
@goStormpath